

Problem Statement: Multi-Agent AI Interview Panel Simulator

What You Need to Build

You will build a team of AI “agents” (like AI teammates) that work together to review a job interview and decide if the candidate should be hired. You will be given a resume, a job description, and an interview transcript — you won’t be told the “correct” answer, because there isn’t one. Your job is to see if your agents can spot the real issues hidden inside.

What You're Given

- 1 Job Description (02_Job_Description.pdf)
- 2 Candidate Resumes (03_Resume_A.pdf, 04_Resume_B.pdf)
- 2 Interview Transcripts, one for each candidate (05_Transcript_A.pdf, 06_Transcript_B.pdf)

What You Must Build

1. A Candidate Profile Builder

A part of your system that reads the resume + transcript and pulls out the basic facts (skills, experience, claims made) — this is the shared information all your agents will use.

2. At least 4 AI Personas (agents), each with a different job

- Technical Agent — checks technical skill and depth
- HR / Culture Agent — checks communication, teamwork, honesty
- Hiring Manager Agent — checks if this person is worth hiring for the role
- Skeptic Agent — looks for contradictions, exaggeration, red flags

Important rule: each agent must first give its opinion on its own, without seeing what the other agents said (there should be a separate LLM call, no agent may access another agent's conclusions before the debate stage). And every opinion must be backed by a real quote or fact from the transcript/resume — not just a random score.

3. A Debate Step

After all agents give their independent opinion, they must talk to each other. At least one agent must directly respond to another agent’s point — agreeing, disagreeing, or changing its own opinion because of it. Just showing 4 opinions side-by-side does NOT count as a debate.

4. A Final Decision Step (not simple averaging)

Your system should not just average the 4 scores to get a final answer. You need some kind of reasoning step where the agents’ evidence and confidence are weighed to reach a final decision — this is the most important part of the challenge.

5. A Final Report

For each candidate, show: final recommendation, confidence level, strengths, concerns, and any disagreement between the agents that wasn’t fully resolved. Integrating a voice debate session between different personas is a bonus point.

Rules

- Every score/opinion must point to a specific quote or fact — no unexplained numbers.
- You must be able to show the moment an agent’s opinion changed after seeing another agent’s point (this proves it’s a real multi-agent system, not one big prompt pretending to be four agents).
- If there isn’t enough information to judge something, say so — don’t make up a score.
- You must process both candidates. Comparing/ranking the two candidates against each other is a bonus, not required.

Judging (100 pts)

- Are the 4 agent personas actually different and independent? — 20 pts
- Quality of the debate + how the final decision is reached — 20 pts
- Can every decision be traced back to evidence? — 15 pts
- How well the system/code is built — 15 pts
- Does it handle unclear or missing info sensibly? — 10 pts
- How easy and clear is it to use? — 10 pts
- Anything creative/extra you add — 10 pts

Tip: most teams spend too much time on the agents and rush the debate step. Don't make that mistake.